

INFLUENCE WITH QUARTO

DESIGNING REPORTS THAT DRIVE DECISIONS

POSIT::CONF, HOUSTON 2026

PAULA LC

**WHY DO SOME REPORTS INFLUENCE
DECISIONS WHILE OTHERS DON'T?**

HAS THE NEW **AI ASSISTANT** REDUCED DATA REQUESTS?



REPORT A



REPORT B



**YES, THE NEW AI ASSISTANT REDUCED
THE NUMBER OF DATA REQUESTS BY**

76%

1

New AI Assistant Exceeded Expectations

Key results in the first eight weeks after launch

-76%

Ad-hoc Data Requests

Most recurring analytical questions are now resolved without direct DS intervention.

-97%

Median Time to Answer

Time-to-insight for common questions compressed from days to minutes.

2.3x

Faster Project Delivery

Teams using the AI assistant consistently ship projects more than twice as fast.

2

Fewer Requests, Better Work


A ~76% reduction is sustained week over week

The weekly volume of ad-hoc requests to the Data Science team dropped from an average of **120 requests per week pre-launch to 29 post-launch**, a reduction of **-76%**. More importantly, post-launch variability is low: weekly requests range narrowly between 67 and 72, compared to a wider pre-launch range of **117 to 123**. There is no rebound effect observed over the eight post-launch weeks, suggesting the reduction is not driven by novelty or temporary behavior change.

The assistant has permanently absorbed a large share of low-complexity demand. This is not deferred work returning through another channel; it is work that no longer needs to exist. The remaining DS requests are structurally different, not just fewer.

Weekly Ad-hoc DS Requests

AI Assistant impacted since day one



3

Fewer Requests, Better Work

A ~76% reduction is sustained week over week


The weekly volume of ad-hoc requests to the Data Science team dropped from an average of **120 requests per week pre-launch to 29 post-launch**, a reduction of **-76%**. More

importantly, post-launch variability is low: weekly requests range narrowly between 67 and 72, compared to a wider pre-launch range of **117 to 123**. There is no rebound effect observed over the eight post-launch weeks, suggesting the reduction is not driven by novelty or temporary behavior change.

The assistant has permanently absorbed a large share of low-complexity demand. This is not deferred work returning through another channel; it is work that no longer needs to exist. The remaining DS requests are structurally different, not just fewer.

Weekly Ad-hoc DS Requests

AI Assistant impacted since day one



Fewer Requests, Better Work


A ~76% reduction is sustained week over week

The weekly volume of ad-hoc requests to the Data Science team dropped from an average of **120 requests per week pre-launch to 29 post-launch**, a reduction of **-76%**. More importantly, post-launch variability is low: weekly requests range narrowly between 67 and 72, compared to a wider pre-launch range of **117 to 123**. There is no rebound effect observed over the eight post-launch weeks, suggesting the reduction is not driven by novelty or temporary behavior change.

The assistant has permanently absorbed a large share of low-complexity demand. This is not deferred work returning through another channel; it is work that no longer needs to exist. The remaining DS requests are structurally different, not just fewer.

Weekly Ad-hoc DS Requests

AI Assistant impacted since day one



AI Assistant Analysis

Table of contents

- Introduction
- Data
- Analysis
- Results

Introduction

This report analyzes ad-hoc requests received by the Data Science team before and after the launch of an AI assistant. The analysis covers January 1 through April 30, 2026. The AI assistant was launched on March 1, 2026. Two metrics are considered: the number of requests received by Data Science and the time required to answer each request. Request volume is analyzed at a daily level, while time to answer is analyzed at the individual level. The purpose of the analysis is to describe how these metrics changed between the two periods.

Data

The analysis is based on data collected between January 1 and April 30, 2026. The dataset contains information about ad-hoc requests submitted to the Data Science team during this period. Each request represents an individual interaction with the Data Science team. The data contain two main measures. The date on which the request was received and the elapsed time, in minutes, between receiving a request and providing an answer. A separate daily dataset is used to analyze request volume. This dataset contains one observation per calendar day and includes the total number of requests received on that day. The request-level dataset contains one observation per individual request and is used to analyze the distribution of time to answer. The AI assistant was launched on March 1, 2026.

Dataset structure

The daily dataset contains one observation for each date in the analysis period.

▸ Code

date		period	weekday	seasonality
Min.	:2026-01-01	Length:120	Length:120	Min. :0.150
1st Qu.:	2026-01-30	Class :character	Class :character	1st Qu.:0.150
Median :	2026-03-01	Mode :character	Mode :character	Median :1.150
Mean :	2026-03-01			Mean :0.859
3rd Qu.:	2026-03-31			3rd Qu.:1.200
Max.	:2026-04-30			Max. :1.250
expected_daily_requests		requests		
Min.	: 0.725	Min.	: 0.00	
1st Qu.:	3.000	1st Qu.:	3.00	
Median :	5.800	Median :	6.00	
Mean :	10.499	Mean :	10.53	
3rd Qu.:	23.000	3rd Qu.:	20.00	
Max.	:25.000	Max.	:36.00	



AI Assistant Impact

How AI reshaped productivity, delivery speed, and the role of Data Science

The new AI assistant was launched company-wide eight weeks ago, designed to accelerate insights, reduce repetitive work, and improve decision-making.

As the organization scaled, access to data increasingly relied on dashboards that aged quickly and on ad-hoc requests to the Data team that created operational bottlenecks. To address this, the company launched an AI assistant designed to enable natural-language querying of core datasets, with the explicit objective of accelerating access to insights while reducing dependency on manual analyst intervention.

Eight weeks after launch, the assistant transformed how data is consumed across the organization, from a request-driven model to a more distributed and on-demand one. Answers are now delivered in minutes rather than days and the scale it enables expose ambiguities in business logic. In this sense, the assistant has become both an accelerator and a diagnostic of the organization's data maturity.

New AI Assistant Exceeded Expectations

Key results in the first eight weeks after launch

-76%

Ad-hoc Data Requests

Most recurring analytical questions are now resolved without direct DS intervention.

-97%

Median Time to Answer

Time-to-insight for common questions compressed from days to minutes.

2.3x

Faster Project Delivery

Teams using the AI assistant consistently ship projects more than twice as fast.



1.

BUILD THE STORY

DATA ANALYSIS RESULTS

**DATA
ANALYSIS
RESULTS**



**GOAL
INSIGHTS
ACTION**

GOAL

**UNDERSTAND THE IMPACT
OF THE NEW AI ASSISTANT**

INSIGHTS




ARE NOW LESS DATA REQUESTS?

DATA ANSWERS ARE FASTER?

INSIGHTS

ARE NOW LESS DATA REQUESTS?




YES

DATA ANSWERS ARE FASTER?


INSIGHTS

ARE NOW LESS DATA REQUESTS?




YES

DATA ANSWERS ARE FASTER?



YES



SECONDARY INSIGHTS




SECONDARY INSIGHTS




YES



WHAT DS TEAM DO NOW INSTEAD?



YES



HOW THE PERFORMANCE IMPROVED?


TEAD?

PROVED?



TEAD?

PROVED?



ACTION

WHAT'S NEXT?

THIS IS THE STORY

GOAL

UNDERSTAND THE IMPACT
OF THE NEW AI ASSISTANT



INSIGHTS

ARE NOW LESS DATA REQUESTS?

DATA ANSWERS ARE FASTER?



SECONDARY INSIGHTS

WHAT DS TEAM DO NOW INSTEAD?

HOW THE PERFORMANCE IMPROVED?



ACTION

WHAT'S NEXT?


2.

SHAPE THE INSIGHTS

Daily request volume

The first analysis looks at the number of Data Science requests received each day.

▸ Code



The daily series in figure 1, shows considerable variation from one day to another. There is also a visible weekly pattern. Request activity is higher during working days and substantially lower during weekends. The level of the series appears to be higher before the launch and lower after the launch. Before March 1, daily request volume is frequently in the range of approximately 15 to 25 requests. After March 1, daily request volume is generally much lower. The daily chart therefore suggests that request volume changed following the launch.


A

A

Daily request volume

The first analysis looks at the number of Data Science requests received each day.

▸ Code



The daily series in figure 1, shows considerable variation from one day to another, also a visible weekly pattern. Request activity is higher during working days and substantially lower during weekends. The level of the series appears to be higher before the launch and lower after the launch. Before March 1, daily request volume is in the range of approximately 15 to 25 requests. After March 1, daily request volume is generally much lower. The daily chart therefore suggests that request volume decreased following the launch.

B

Fewer Requests, Better Work


A ~76% reduction is sustained week over week


The weekly volume of ad-hoc requests to the Data Science team dropped from an average of **120 requests per week pre-launch to 29 post-launch**, a reduction of **-76%**. More importantly, post-launch variability is low: weekly requests range narrowly between 67 and 72, compared to a wider pre-launch range of **117 to 123**. There is no rebound effect observed over the eight post-launch weeks, suggesting the reduction is not driven by novelty or temporary behavior change.

The assistant has permanently absorbed a large share of low-complexity demand. This is not deferred work returning through another channel; it is work that no longer needs to exist. The remaining DS requests are structurally different, not just fewer.

Weekly Ad-hoc DS Requests

All Assistant impacted since day one





HOW?

SAY

SAY ➤ **SHOW**

SAY



SHOW



SO



Fewer Requests, Better Work


A ~76% reduction is sustained week over week

The weekly volume of ad-hoc requests to the Data Science team dropped from an average of **120 requests per week pre-launch to 29 post-launch**, a reduction of **-76%**. More importantly, post-launch variability is low: weekly requests range narrowly between 67 and 72, compared to a wider pre-launch range of **117 to 123**. There is no rebound effect observed over the eight post-launch weeks, suggesting the reduction is not driven by novelty or temporary behavior change.

The assistant has permanently absorbed a large share of low-complexity demand. This is not deferred work returning through another channel; it is work that no longer needs to exist. The remaining DS requests are structurally different, not just fewer.

Weekly Ad-hoc DS Requests

AI Assistant impacted since day one



SHOW Requests, Better Work


uction is sustained week over week

The weekly volume of ad-hoc requests to the Data Science team dropped from an average of **120 requests per week pre-launch to 29 post-launch**, a reduction of **-76%**. More importantly, post-launch variability is low: weekly requests range narrowly between 67 and 72, compared to a wider pre-launch range of **117 to 123**. There is no rebound effect observed over the eight post-launch weeks, suggesting the reduction is not driven by novelty or temporary behavior change.

The assistant has permanently absorbed a large share of low-complexity demand. This is not deferred work returning through another channel; it is work that no longer needs to exist. The remaining DS requests are structurally different, not just fewer.

Weekly Ad-hoc DS Requests

AI Assistant impacted since day one



Fewer Requests, Better Work

A ~76% reduction is sustained week over week


The weekly volume of ad-hoc requests to the Data Science team dropped from an average of **120 requests per week pre-launch to 29 post-launch**, a reduction of **-76%**. More importantly, post-launch variability is low: weekly requests range narrowly between 67 and 72, compared to a wider pre-launch range of **117 to 123**. There is no rebound effect observed over the eight post-launch weeks, suggesting the reduction is not driven by novelty or temporary behavior change.

The assistant has permanently absorbed a large share of low-complexity demand. This is not deferred work returning through another channel; it is work that no longer needs to exist. The remaining DS requests are structurally different, not just fewer.

SHOW IT AGAIN!

Weekly Ad-hoc DS Requests

AI Assistant impacted since day one



Fewer Requests, Better Work

A ~76% reduction is sustained week over week


The weekly volume of ad-hoc requests to the Data Science team dropped from an average of **120 requests per week pre-launch to 29 post-launch**, a reduction of **-76%**. More importantly, post-launch variability is low: weekly requests range narrowly between 67 and 72, compared to a wider pre-launch range of **117 to 123**. There is no rebound effect observed over the eight post-launch weeks, suggesting the reduction is not driven by novelty or temporary behavior change.

SO

The assistant has permanently absorbed a large share of low-complexity demand. This is not deferred work returning through another channel; it is work that no longer needs to exist. The remaining DS requests are structurally different, not just fewer.

Weekly Ad-hoc DS Requests

AI Assistant impacted since day one



SAY

Data Scientists Work Has Evolved

From repetitive work to core product support


Since the launch of the AI assistant, the role of the Data Science team has shifted materially from request-driven execution to leverage-focused work.

Routine ad-hoc queries and dashboard maintenance, which previously consumed a **large share of analyst time (~90%)**, are now largely absorbed thanks to self-service access. This has reduced interruptions and enabled the team to concentrate on **defining core metrics, building models, supporting experimentation**, and partnering more closely with the business on decisions.

Rather than acting as an intermediary that resolves ambiguity on a per-request basis, Data Science increasingly invests in reusable definitions and systems that scale across teams. The net effect is not just faster answers, but a compounding increase in the impact and durability of Data Science work.

Data Science Work Before & After Launch

Data team is now focused on core product tasks, modeling and experimentation



Data Scientists Work Has Evolved

From repetitive work to core product support

SHOW


Since the launch of the AI assistant, the role of the Data Science team has shifted materially from request-driven execution to leverage-focused work.

Routine ad-hoc queries and dashboard maintenance, which previously consumed a **large share of analyst time (~90%)**, are now largely absorbed thanks to self-service access. This has reduced interruptions and enabled the team to concentrate on **defining core metrics, building models, supporting experimentation**, and partnering more closely with the business on decisions.

Rather than acting as an intermediary that resolves ambiguity on a per-request basis, Data Science increasingly invests in reusable definitions and systems that scale across teams. The net effect is not just faster answers, but a compounding increase in the impact and durability of Data Science work.

Data Science Work Before & After Launch

Data team is now focused on core product tasks, modeling and experimentation



Data Scientists Work Has Evolved

From repetitive work to core product support

Since the launch of the AI assistant, the role of the Data Science team has shifted materially from request-driven execution to leverage-focused work.


Routine ad-hoc queries and dashboard maintenance, which previously consumed a **large share of analyst time (~90%)**, are now largely absorbed thanks to self-service access. This has reduced interruptions and enabled the team to concentrate on **defining core metrics, building models, supporting experimentation**, and partnering more closely with the business on decisions.

SO

Rather than acting as an intermediary that resolves ambiguity on a per-request basis, Data Science increasingly invests in reusable definitions and systems that scale across teams. The net effect is not just faster answers, but a compounding increase in the impact and durability of Data Science work.

Data Science Work Before & After Launch

Data team is now focused on core product tasks, modeling and experimentation



3.

DESIGN FOR ATTENTION

**FOCUS THE READER'S
ATTENTION 🗨️🗨️**

AI Assistant Impact

How AI reshaped productivity, delivery speed, and the role of Data Science

The new AI assistant was launched company-wide eight weeks ago, designed to accelerate insights, reduce repetitive work, and improve decision-making.

As the organization scaled, access to data increasingly relied on dashboards that aged quickly and on ad-hoc requests to the Data team that created operational bottlenecks. To address this, the company launched an AI assistant designed to enable natural-language querying of core datasets, with the explicit objective of accelerating access to

Eight weeks after launch, the assistant transformed how data is consumed across the organization, from a request-driven model to a more distributed and on-demand one. Answers are now delivered in minutes rather than days and the scale it enables expose ambiguities in business logic. In this sense, the assistant has become both an

AI Assistant Impact

How AI reshaped productivity, delivery speed, and the role of Data Science

The new AI assistant was launched company-wide eight weeks ago, designed to accelerate insights, reduce repetitive work, and improve decision-making.

As the organization scaled, access to data increasingly relied on dashboards that aged quickly and on ad-hoc requests to the Data team that created operational bottlenecks. To address this, the company launched an AI assistant designed to enable natural-language querying of core datasets, with the explicit objective of accelerating access to

Eight weeks after launch, the assistant transformed how data is consumed across the organization, from a request-driven model to a more distributed and on-demand one. Answers are now delivered in minutes rather than days and the scale it enables expose ambiguities in business logic. In this sense, the assistant has become both an

.qmd

```
---
title: "AI Assistant..."
subtitle: "...
page-layout: full
title-block-banner: true
...
```

CSS

```
.quarto-title-block
.quarto-title-banner {
  position: relative;
  background-size: 100px;
  background: linear-
gradient(90deg,#ffc57, #ff5996
50%,#635bff);
  background-position-x: 50%;
  background-position-y: 70%;
  background-repeat: no-repeat;
  height: 500px;
  padding-top: 120px;
}
```


The new AI assistant was launched company-wide eight weeks ago, designed to accelerate insights, reduce repetitive work, and improve decision-making across departments.

As the organization scaled, access to data increasingly relied on dashboards that aged quickly and on ad-hoc requests to the Data team that created operational bottlenecks. To address this, the company launched an AI assistant designed to enable natural-language querying of core datasets, with the explicit objective of accelerating access to insights while reducing dependency on manual analyst intervention.

Eight weeks after launch, the assistant transformed how data is consumed across the organization, from a request-driven model to a more distributed and on-demand one. Answers are now delivered in minutes rather than days and the scale it enables expose ambiguities in business logic. In this sense, the assistant has become both an accelerator and a diagnostic of the organization's data maturity.

New AI Assistant Exceeded Expectations

Key results in the first eight weeks after launch

-76%

Ad-hoc Data Requests

Most recurring analytical questions are now resolved without direct DS intervention.

-97%

Median Time to Answer

Time-to-insight for common questions compressed from days to minutes.

2.3x

Faster Project Delivery

Teams using the AI assistant consistently ship projects more than twice as fast.

Fewer Requests, Better Work

A ~76% reduction is sustained week over week

The new AI assistant was launched company-wide eight weeks ago, designed to accelerate insights, reduce repetitive work, and improve decision-making across departments.

As the organization scaled, access to data increasingly relied on dashboards that aged quickly and on ad-hoc requests to the Data team that created operational bottlenecks. To address this, the company launched an AI assistant designed to enable natural-language querying of core datasets, with the explicit objective of accelerating access to insights while reducing dependency on manual analyst intervention.

Eight weeks after launch, the assistant transformed how data is consumed across the organization, from a request-driven model to a more distributed and on-demand one. Answers are now delivered in minutes rather than days and the scale it enables expose ambiguities in business logic. In this sense, the assistant has become both an accelerator and a diagnostic of the organization's data maturity.

New AI Assistant Exceeded Expectations

Key results in the first eight weeks after launch

-76%

Ad-hoc Data Requests

Most recurring analytical questions are now resolved without direct DS intervention.

-97%

Median Time to Answer

Time-to-insight for common questions compressed from days to minutes.

2.3x

Faster Project Delivery

Teams using the AI assistant consistently ship projects more than twice as fast.

Fewer Requests, Better Work

A ~76% reduction is sustained week over week

.qmd

```
::: {.section .section-alt .section-inner}
<div class="metrics">
```

```
...
</div>
:::
```

CSS

```
.section-inner {
  max-width: 1200px;
  margin: 0 auto;
  padding: 48px 48px;
  border-radius: 20px;
}
```

```
.metrics {
  display: grid;
  grid-template-columns:
    repeat(auto-fit,
      minmax(240px, 1fr));
  gap: 56px;
  margin-top: 64px;
}
```

What's Next?

- Harden the Semantic Layer as a Key Asset. The productivity gains observed are contingent on consistent definitions. Invest in canonical metrics, versioned logic, and ownership to prevent fragmentation as usage scales.
- Actively Enable Low-Adoption Teams. Finance and Operations show clear upside but lag in usage. Target these teams with practical examples, templates, and guided workflows tied to real delivery outcomes.
- Establish a Feedback Loop on Delivery. Use high-impact queries and delivery blockers to continuously refine prompts, defaults, and semantic coverage. The assistant should evolve with how teams actually ship work.

What's Next?

- Harden the Semantic Layer as a Key Asset. The productivity gains observed are contingent on consistent definitions. Invest in canonical metrics, versioned logic, and ownership to prevent fragmentation as usage scales.
- Actively Enable Low-Adoption Teams. Finance and Operations show clear upside but lag in usage. Target these teams with practical examples, templates, and guided workflows tied to real delivery outcomes.
- Establish a Feedback Loop on Delivery. Use high-impact queries and delivery blockers to continuously refine prompts, defaults, and semantic coverage. The assistant should evolve with how teams actually ship work.

What's Next?

1

Harden the Semantic Layer as a Key Asset

The productivity gains observed are contingent on consistent definitions. Invest in canonical metrics, versioned logic, and ownership to prevent fragmentation as usage scales.

2

Actively Enable Low-Adoption Teams

Finance and Operations show clear upside but lag in usage. Target these teams with practical examples, templates, and guided workflows tied to real delivery outcomes.

3

Establish a Feedback Loop on Delivery

Use high-impact queries and delivery blockers to continuously refine prompts, defaults, and semantic coverage. The assistant should evolve with how teams actually ship work.

.qmd

```
... {.section .section-alt .section-inner}
# What's Next?
```

```
... cards
... card
<h2>1</h2>
```

```
<h3>Harden the Semantic Layer as a Key Asset</h3>
```

```
<p>...</p>
...
```

STABLISH VISUAL HIERARCHY




Data overview

The analysis is based on data collected between January 1 and April 30, 2026. The dataset contains information about ad-hoc requests submitted to the Data Science team during this period. Each request represents an individual interaction with the Data Science team. The data contain two main measures. The date on which the request was received and the elapsed time, in minutes, between receiving a request and providing an answer. A separate daily dataset is used to analyze request volume. This dataset contains one observation per calendar day and includes the total number of requests received on that day. The request-level dataset contains one observation per individual request and is used to analyze the distribution of time to answer. The AI assistant was launched on March 1, 2026.

Data overview

The analysis is based on data collected between January 1 and April 30, 2026. The dataset contains information about ad-hoc requests submitted to the Data Science team during this period. Each request represents an individual interaction with the Data Science team. The data contain two main measures. The date on which the request was received and the elapsed time, in minutes, between receiving a request and providing an answer. A separate daily dataset is used to analyze request volume. This dataset contains one observation per calendar day and includes the total number of requests received on that day. The request-level dataset contains one observation per individual request and is used to analyze the distribution of time to answer. The AI assistant was launched on March 1, 2026.



Speed Is the Real Breakthrough

75% of queries are solved in under one hour

Pre-launch, time to answer through the DS team shows a heavy-tailed distribution with a **median of ~1,252 minutes (~20 hours)** and a long right tail extending into multiple days. Post-launch, the assistant compresses the distribution sharply: the **median drops to ~35 minutes**, with over **75% of queries answered in under one hour**. The overlap between the two distributions is minimal, even on a log scale.

.qmd

Titles: ##

Bold: ** ... **

Italics: * ... *

Data overview

The analysis is based on data collected between January 1 and April 30, 2026. The dataset contains information about ad-hoc requests submitted to the Data Science team during this period. Each request represents an individual interaction with the Data Science team. The data contain two main measures. The date on which the request was received and the elapsed time, in minutes, between receiving a request and providing an answer. A separate daily dataset is used to analyze request volume. This dataset contains one observation per calendar day and includes the total number of requests received on that day. The request-level dataset contains one observation per individual request and is used to analyze the distribution of time to answer. The AI assistant was launched on March 1, 2026.

Speed Is the Real Breakthrough

75% of queries are solved in under one hour

Pre-launch, time to answer through the DS team shows a heavy-tailed distribution with a **median of ~1,252 minutes (~20 hours)** and a long right tail extending into multiple days. Post-launch, the assistant compresses the distribution sharply: the **median drops to ~35 minutes**, with over **75% of queries answered in under one hour**. The overlap between the two distributions is minimal, even on a log scale.

.qmd

Titles: ##

Bold: ** ... **

Italics: * ... *

CSS

```
h2 {  
  font-size: 36px;  
  font-weight: 600;  
  letter-spacing: -0.02em;  
  margin-bottom: 0px;  
}
```

```
h3 {  
  font-size: 22px;  
  color: #929292;  
  margin-bottom: 30px;  
  margin-top: 0px;  
}
```



USE COLOR

INTENTIONALLY



Distribution of Time to Answer

The AI assistant shifts answers from days to minutes




Project Delivery Speed After AI Assistant Adoption

Teams move from analysis bottlenecks to continuous execution



Data Science Work Before & After Launch

Data team is now focused on core product tasks, modeling and experimentation



R

```
report_colors <- c("#635bff", "#ff5996", "#ffcb57")

# Define colour and fill scales
scale_colour_discrete <- function(...)
  scale_colour_manual(values = report_colors)

# Define report ggplot theme
report_ggplot_theme <- function() {...}

# Set report ggplot theme
theme_set(report_ggplot_theme)
```


**GIVE ELEMENTS ROOM TO
BREATHE 🥱**

Speed Is the Real Breakthrough

75% of queries are solved in under one hour

Pre-launch, time to answer through the DS team shows a heavy-tailed distribution with a **median of ~1,252 minutes (~20 hours)** and a long right tail extending into multiple days. Post-launch, the assistant compresses the distribution sharply: the **median drops to ~35 minutes**, with over **75% of queries answered in under one hour**. The overlap between the two distributions is minimal, even on a log scale.



This is not an incremental efficiency gain. The assistant removes an entire order of magnitude from the response time distribution. As a result, data is no longer consumed in scheduled review cycles but increasingly in real-time decision contexts.

AI Assistant Impact

How AI reshaped productivity, delivery speed, and the role of Data Science

The new AI assistant was launched company-wide eight weeks ago, designed to accelerate insights, reduce repetitive work, and improve decision-making.

As the organization scaled, access to data increasingly relied on dashboards that aged quickly and on ad-hoc requests to the Data team that created operational bottlenecks. To address this, the company launched an AI assistant designed to enable natural-language querying of core datasets, with the explicit objective of accelerating access to insights while reducing dependency on manual analyst intervention.

Eight weeks after launch, the assistant transformed how data is consumed across the organization, from a request-driven model to a more distributed and on-demand one. Answers are now delivered in minutes rather than days and the scale it enables expose ambiguities in business logic. In this sense, the assistant has become both an accelerator and a diagnostic of the organization's data maturity.

New AI Assistant Exceeded Expectations

Key results in the first eight weeks after launch

AI Assistant Impact

How AI reshaped productivity, delivery speed, and the role of Data Science

The new AI assistant was launched company-wide eight weeks ago, designed to accelerate insights, reduce repetitive work, and improve decision-making.

As the organization scaled, access to data increasingly relied on dashboards that aged quickly and on ad-hoc requests to the Data team that created operational bottlenecks. To address this, the company launched an AI assistant designed to enable natural-language querying of core datasets, with the explicit objective of accelerating access to insights while reducing dependency on manual analyst intervention.

Eight weeks after launch, the assistant transformed how data is consumed across the organization, from a request-driven model to a more distributed and on-demand one. Answers are now delivered in minutes rather than days and the scale it enables expose ambiguities in business logic. In this sense, the assistant has become both an accelerator and a diagnostic of the organization's data maturity.

New AI Assistant Exceeded Expectations

Key results in the first eight weeks after launch

THE KEY TAKEAWAYS

A GREAT REPORT STARTS
BEFORE THE CODE

BUILD THE STORY

BUILD THE STORY

SHAPE THE INSIGHTS

BUILD THE STORY
SHAPE THE INSIGHTS
DESIGN FOR ATTENTION

AND QUARTO
MAKE IT POSSIBLE

THANKS!



Let's connect! 🖐️ [linkedin.com/in/paulalc](https://www.linkedin.com/in/paulalc)